

Applied Machine Learning — Summer 2026

Task 1: sentiment classification of movie-review snippets contaminated with spam email. **Task 2:** 5-point face-landmark alignment. All experiments ran on an Apple M4 Mac Mini, deep models on PyTorch/MPS. Every number and figure below comes from the code in this archive, whose `tests_sanity.py` files verify the mechanics behind them.

Task 1 — Sentiment analysis

1.1 The problem

The training set is 11,503 documents labelled 0 (negative) or 1 (positive), but roughly a quarter are not reviews: they are Enron-style emails injected as spam. The brief requires spam in the test set to get a label that is neither 0 nor 1, so this is a three-way problem with labels for only two classes; I use -1 as the dummy.

That shapes the architecture (Figure 1). Rather than asking one classifier to learn "spam" from labels that never mark it, two decoupled stages: an unsupervised structural gate deciding *whether this is a review*, and a sentiment model that only sees what survives it. They share no features, so a failure in one cannot corrupt the other.




Figure 1: Task 1 pipeline. The upper branch decides genre from surface structure, the lower sentiment from lexical content; they are deliberately disjoint.

1.2 Preprocessing and representation

Spam gate. Ten surface statistics per document (`src/structure_features.py`): log length, line-break, digit, non-alphanumeric and uppercase ratios, mean token length, and binary flags for RFC-822 headers, URLs and email addresses. These are genre cues, not topic cues — deliberately, because a gate built on vocabulary stops working the moment the spam changes subject. A two-component Gaussian mixture is fitted to the standardised features with no labels at all.

Sentiment. A union of word (1–2 gram) and character (`char_wb`, 3–5 gram) TF-IDF, sublinear term frequency, `min_df=3`. Character n-grams matter because the snippets are short and noisy, capturing *wasn't*, sub-word affixes and typos without stemming.

My hypothesis was that with character n-grams present, hand-engineered preprocessing would be redundant. Table 1 tests it with the classifier fixed, so each cell isolates one decision.

Preprocessing	word (1,2)	char (3,5)	word + char
lowercase only	0.741	0.749	0.758
+ stopword removal	0.767	0.750	0.758
+ lemmatisation	0.743	0.765	0.767
+ negation marking	0.770	0.769	0.780


Table 1: Preprocessing $\times$ representation ablation, validation accuracy on genuine reviews.

Half the hypothesis survived. Stopword removal and lemmatisation buy nothing once character n-grams are present — lemmatisation even *costs* the word-only model accuracy, destroying the distinctions separating *disappoints* from *disappointed*. Negation marking is the exception, adding 2.2 points: character n-grams see that a negator is *present* but not what it applies to, and a bag of n-grams cannot represent scope. I adopted it as the deployed preprocessor — a decision made by measurement, against my prior.

1.3 Prediction methods

Supervised binary classification with a routed third label. Four models spanning different inductive biases:

1. **Word-list classifier** — the interpretable floor: the 400 words most over-represented in each class by document frequency, predicting positive when positive hits exceed negative.
2. **Multinomial naive Bayes** on the TF-IDF union ($\alpha = 0.3$). Generative, closed-form.
3. **Calibrated linear SVM** (LinearSVC, $C = 1$, `class_weight="balanced"`) on the same union. Hinge loss, max-margin; Platt scaling (sigmoid, 3-fold) supplies the routing probabilities.
4. **BiLSTM over pretrained GloVe embeddings** (Figure 2) — the structurally different second method. Every model above is a bag of substrings that cannot represent word order, yet the mined failures are overwhelmingly compositional. Token ids padded to 60, frozen GloVe-100d table (90.4% coverage), one bidirectional LSTM layer of 128 units per direction, max-over-time pooling, dropout 0.4, cross-entropy loss, Adam $1e-3$.



The two prediction methods; their probabilities are averaged in the deployed soft-vote ensemble

Figure 2: The two prediction heads, whose probabilities are averaged in the deployed system.

Hyperparameters, and why. Embeddings frozen: 8.5k short sentences cannot re-estimate $400k \times 100$ parameters without memorising them, and freezing preserves the geometry GloVe learned from 6B tokens. One layer, not two — the sentences average 21 tokens, so depth adds capacity where there is no signal to fit. Max-over-time pooling rather than the final hidden state, so a decisive clause anywhere can carry the prediction. Vocabulary pruning is switched *off* (`min_count=1`, worth 1.5 points) because GloVe already supplies a good vector for a word seen once. Early stopping runs against a 10% slice held out of the *training* set, never validation — otherwise the BiLSTM row would be selected on the data it is scored on.

1.4 Quantitative results

Model	Val. accuracy	Val. F1	Train time (s)
Soft-vote ensemble	0.811	0.805	49.1
Calibrated linear SVM	0.790	0.792	1.2
Multinomial naive Bayes	0.779	0.776	0.8
BiLSTM + GloVe	0.774	0.764	47.1
Word-list floor	0.603	0.695	0.05

Table 2: Validation accuracy on the 1,066 genuine reviews. Wall-clock on the M4.

The BiLSTM does **not** beat the sparse models alone: a sequence model is the wrong shape for 8.5k single-sentence examples where lexical cues carry most of the signal. All three non-trivial models land within 1.6 points of each other, suggesting a data ceiling rather than a modelling failure — Section 1.5 shows label noise consistent with that.

What it *does* provide is decorrelated errors, being the only member that can see word order. Averaging the three calibrated probability vectors — free, since nothing is refitted — reaches 0.811, beating every member by 2.1 points. That is the payoff for a second method, and it exists only because the method was structurally different rather than another linear model on the same features.

Including spam routing, three-way validation accuracy is **0.856**. Figure 3 shows the spam class perfectly separated (331/331, no leakage either way) and residual sentiment errors near-symmetric: 82 positives called negative against 119 the reverse.




Figure 3: Three-way validation confusion matrix, deployed system.

1.5 Qualitative analysis

The 12 highest-confidence errors are in `results/failure_cases.csv`; four patterns explain nearly all of them.

"a lightweight, uneven action comedy that freely mingles french, japanese and hollywood cultures." — gold 0, predicted 1 at 0.96.

Damning with faint praise. Every content word is neutral-to-positive; the negative verdict rests entirely on *lightweight* and *uneven*. The model cannot represent understatement.

"neither the funniest film eddie murphy nor robert de niro has ever made, showtime is nevertheless efficiently amusing." — gold 1, predicted 0.

Mixed polarity with a concessive pivot. Negative for two-thirds of its length, reversing on *nevertheless*. A bag of n-grams sums both halves and the negative half is longer; negation marking cannot help, the reversal being discourse-level, not syntactic.

"you're never quite sure where self-promotion ends and the truth begins. but ... you're too interested to care." — gold 1, predicted 0.

Negation of a negative. Both phrases are positive in effect and negative in surface form — exactly the compositional case the BiLSTM was added for, and here it is right where the SVM is wrong.

"idiotic and ugly." — gold 1, predicted 0 at 0.92 confidence.

Label noise. Not a model failure. Several such items exist, placing a hard ceiling on achievable accuracy and explaining the ~0.79 plateau.

1.6 Auditing the spam gate

The gate scores precision = recall = F1 = 1.000 against a high-precision RFC-822 header rule. That is close to circular: both read document structure, so their agreement is not independent evidence. Two things that *are*:

A hand audit. 50 randomly sampled gated documents, checked against evidence the gate cannot see — business and email *vocabulary*, where every gate feature is a surface statistic with no access to the lexicon. 44 of 50 (88%) contain it; the other six are short forwarded fragments with no distinctive words either way, not misfires on reviews. The sample is in `results/spam_gate_audit.csv` with excerpts, so the claim can be checked rather than trusted.

A zero false-positive rate out of domain. Over NLTK `movie_reviews` — genuine reviews, no spam, far longer than anything in training — the gate fires **0 times in 2,000**. Since length is its strongest feature, that is the sharpest available test of its false-positive mode.

1.7 External generalisation

On NLTK `movie_reviews` (2,000 full-length critiques) the deployed system scores **0.739** against 0.811 in domain. The shape of the loss is the interesting part: recall is 0.979 on the negative class but 0.498 on the positive — a near-unanimous negative predictor.

This is a length/domain shift, not a defect in the model (Figure 4). Training documents average 21 words; NLTK documents average 746. A long document accumulates far more weighted terms, and because the training vocabulary's discriminative mass skews negative, the score drifts down with length: mean P(positive) falls from 0.491 to 0.333, declining across length deciles and crossing below the 0.5 threshold.

What moved is the operating point, not the discrimination; recalibrating would recover much of the gap but needs target labels.

Why external accuracy drops: a length/domain shift




Figure 4: Scores shift negative out of domain (left) and fall with document length (right), explaining the asymmetric recall in Figure 5.




Figure 5: NLTK movie_reviews confusion — the negative-class bias made concrete.

1.8 Test predictions

Written by the worksheet's `save_as_csv`, copied verbatim into `src/submission.py`: no header, 1,434 rows, test-set order untouched. The distribution is 564 negative, 502 positive and 368 spam-dummy (-1), consistent with the 25.8% spam rate seen in training.

Task 2 — Face alignment

2.1 Preprocessing and representation

2,600 training images, 211 validation and 554 test: 256×256 RGB with five landmarks (eyes, nose, mouth corners). Preprocessing: RGB→greyscale (landmark cues are structural, not chromatic), resize to 64×64 with `INTER_AREA`, intensity to [0,1]. The detail that matters is that landmarks are multiplied by the same

(sx, sy) as the image and inverted again before submission; `tests_sanity.py` asserts that round trip is exact. CLAHE equalisation is implemented but left off — it amplifies noise in the dark frames, of which this set has several.

Two representations follow, one per model family: a global HOG descriptor (6×6 cells, 3×3 block L2-Hys normalisation, 9 orientations), and the raw 64×64 intensity map for the CNN.

2.2 Prediction methods

Supervised regression to ten continuous coordinates. Three approaches plus a floor:

1. **Mean face** — the training mean shape for every image. The honest floor.
2. **PCA shape model + HOG ridge regression.** PCA on the 10-D landmark vectors gives a mean shape plus 6 modes explaining 97.7% of shape variance; ridge regression maps the HOG descriptor to those shape *parameters* rather than raw coordinates, constraining predictions to the manifold of plausible faces — the model cannot output a geometrically impossible one. Loss: squared error with L2 penalty ($\alpha = 1$), closed form.
3. **Heatmap-regression CNN with a soft-argmax decode** (Figure 6) — the headline model. Rather than emitting ten numbers from a dense head it predicts one 64×64 Gaussian heatmap per landmark, decoded by spatial expectation; arg-max is not differentiable, which is what heatmap targets and soft-argmax exist to solve. Loss: pixel-wise MSE against the targets plus $0.1 \times$ MSE on the decoded coordinates. Trained twice, with and without augmentation, because the difference is the robustness argument.




Figure 6: Encoder–decoder heatmap CNN, 100,309 parameters. The 16×16 bottleneck supplies the receptive field to place a landmark from global context; the decoder restores resolution.

Hyperparameters, and why. $\sigma = 1.5$ px makes each target blob ~ 7 px across: wide enough to give gradient everywhere near the landmark, tight enough that the expectation is not dragged by a neighbouring blob. The coordinate term is weighted 0.1 in *grid-relative* units — in raw pixels it starts near 10^3 against a heatmap MSE of 10^{-2} , so no weighting keeps the heatmap term relevant and the model degenerates into the direct regressor this design replaces; targets peak at 1.0 rather than summing to 1 for the same reason. Adam at $1e-3$ with cosine annealing, batch 32. Training stops early on validation **AUC-CED**, not training loss, which is dominated by the heatmap term and improves after the decoded coordinates have stopped. Augmentation (flip with the mandatory $0 \leftrightarrow 1$, $3 \leftrightarrow 4$ index swap; $\pm 25^\circ$ rotation; 0.85–1.15 scale; $\pm 8\%$ translation; photometric jitter and noise) expands training 4×; test-time augmentation averages a flipped pass.

Compute. Apple M4 Mac Mini, PyTorch MPS. Shape model 1.2 s on CPU; CNN without augmentation 3.5 min (early-stopped at epoch 76, best 51), with augmentation 9.1 min (stopped at 50, best 25) — fewer

epochs because each shows it four times the data.

2.3 Quantitative results

Two metrics, answering different questions. The **raw Euclidean pixel error at 256×256**, from the worksheet's `euclid_dist`, is what the graders measure. The **inter-ocular-normalised error** — the same distance over the eye-to-eye distance — is scale-invariant and so the fair basis for comparison; AUC-CED is its area under the CED up to 0.10.

Approach	Mean px	Median px	AUC-CED	≤5 px	≤8 px	≤12 px
Mean face (floor)	12.12	10.48	0.182	4.3%	21.8%	51.2%
PCA shape + HOG ridge	6.20	5.36	0.418	34.1%	79.6%	98.1%
Heatmap CNN, no augmentation	7.73	4.69	0.452	41.7%	69.2%	83.4%
Heatmap CNN, augmented	5.59	4.45	0.512	57.3%	86.7%	96.2%

Table 3: Validation results ($n = 211$). Percentages are the fraction of images whose mean per-landmark error is below the threshold — the graded quantity.

The threshold columns are why selection uses them rather than the mean. The non-augmented CNN has a **better** median than the shape model (4.69 vs 5.36 px) but a **worse** mean (7.73 vs 6.20) — a fat tail of failures dragging the average. It wins at 5 px yet loses at 8 px (69.2% vs 79.6%). Only the augmented CNN wins on every column, so it is deployed.




Figure 7: CED in original-resolution pixels, graded thresholds marked. The crossing curves are the fat tail made visible.




Figure 8: Per-landmark error.

For the deployed model the left eye is most accurate (3.89 px) and the right least (7.13 px), mouth corners between (4.58 and 6.36 px). The asymmetry is a dataset property, not an artefact: here the "right" eye is more often the far eye under yaw, so more often foreshortened or occluded — which Section 2.4 confirms.

2.4 Qualitative results

Best validation cases (cnn_aug): predicted (+) vs ground truth (x)




Figure 9: Best four cases (1.7–2.4 px): frontal, evenly lit, unoccluded.

Worst validation cases (cnn_aug): predicted (+) vs ground truth (x)




Figure 10: Worst four cases (18.9–41.3 px), each a different failure mode.

The failures are not random; each has a nameable cause. #104: a hand occludes the lower face, so the mouth landmarks are placed by shape prior alone and land on the hand. #64: a surgical mask removes every mouth cue and the model puts the corners on the mask's edge, the strongest remaining contour. #106: a stylised film poster, not a photograph, with intensity statistics outside the training distribution. #87: near-profile behind heavy glasses.

Occlusion and out-of-distribution imagery dominate, and **the systematic bias is yaw**. Correlating per-image error against pose proxies read off the ground truth gives $r = +0.37$ against $|\text{yaw}|$ but only $+0.10$ against $|\text{roll}|$ (Figure 11). That asymmetry is itself evidence the augmentation worked: rotation augmentation covers roll, so roll costs almost nothing, while nothing in the set synthesises out-of-plane rotation, so yaw remains uncorrected. Fixing it needs 3-D-aware augmentation or yaw-stratified sampling — neither available from a 2-D affine warp.




Figure 11: Error against head pose: roll is flat, yaw is not.

2.5 Robustness

Both CNNs were stressed under increasing corruption, with the ground truth transformed identically in the geometric case (Figure 12).




Figure 12: Robustness to rotation; the gap between the CNNs is the augmentation effect.

Rotation. The augmented model retains 73% of its AUC-CED at 30° ($0.569 \rightarrow 0.415$); the non-augmented one retains 39% ($0.520 \rightarrow 0.202$), so the augmented model at 30° still beats the non-augmented one *undisturbed*. This is the clean demonstration that augmentation teaches the transformation rather than merely regularising: the network saw rotated faces with correctly rotated targets, so what it learned is approximately equivariant. The shape model degrades faster than either ($0.418 \rightarrow 0.162$) despite HOG's

small-rotation tolerance — its PCA basis was fitted on upright faces, so a rotated face is off the shape manifold however stable the descriptor.

Noise. The ordering reverses, which is the more interesting result. Both CNNs collapse by $\sigma = 0.10$ (to 0.030 and 0.013) while the shape model retains 45% of its AUC-CED at $\sigma = 0.20$. HOG's block-wise contrast normalisation divides out precisely the perturbation noise introduces; the CNN's first convolution passes it straight through. The augmented CNN beats the non-augmented one at every noise level (0.268 vs 0.148 at $\sigma = 0.05$), but its training-time noise was $\sigma = 0.02$, so that advantage does not extrapolate. Heatmap regression buys *geometric* robustness, not photometric — a deployment facing sensor noise should train with matched noise or keep the classical model as a fallback.

2.6 Test predictions

The deployed model was run over all 554 test images in the given order, scaled back to 256×256 and written by the worksheet's `save_as_csv` verbatim. All coordinates lie in [62.3, 196.8] — inside the frame, confirming the inverse scaling ran. Figure 13 is the pre-shipping check.




Figure 13: First four test images with predicted landmarks at original resolution.

Use of generative AI

Generative AI (Claude) was used as a coding assistant: scaffolding plotting code, suggesting refactors, and drafting prose that I then edited. All modelling decisions, hyperparameter justifications, experimental design and analysis are my own, and every number reported here was produced by running the code in this archive.

References

- Cootes, T., Taylor, C., Cooper, D. and Graham, J. (1995) 'Active shape models — their training and application', *Computer Vision and Image Understanding*, 61(1), pp. 38–59.
- Dalal, N. and Triggs, B. (2005) 'Histograms of oriented gradients for human detection', *CVPR*, pp. 886–893.
- Das, S. and Chen, M. (2001) 'Yahoo! for Amazon: extracting market sentiment from stock message boards', *Asia Pacific Finance Association Annual Conference*.
- Newell, A., Yang, K. and Deng, J. (2016) 'Stacked hourglass networks for human pose estimation', *ECCV*, pp. 483–499.
- Pang, B., Lee, L. and Vaithyanathan, S. (2002) 'Thumbs up? Sentiment classification using machine learning techniques', *EMNLP*, pp. 79–86.
- Pedregosa, F. et al. (2011) 'Scikit-learn: machine learning in Python', *Journal of Machine Learning Research*, 12, pp. 2825–2830.

- Pennington, J., Socher, R. and Manning, C. (2014) 'GloVe: global vectors for word representation', *EMNLP*, pp. 1532–1543.
- Platt, J. (1999) 'Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods', *Advances in Large Margin Classifiers*, pp. 61–74.
- Salton, G. and Buckley, C. (1988) 'Term-weighting approaches in automatic text retrieval', *Information Processing & Management*, 24(5), pp. 513–523.
- Sun, X., Xiao, B., Wei, F., Liang, S. and Wei, Y. (2018) 'Integral human pose regression', *ECCV*, pp. 529–545.
- Tompson, J., Jain, A., LeCun, Y. and Bregler, C. (2014) 'Joint training of a convolutional network and a graphical model for human pose estimation', *NeurIPS*, pp. 1799–1807.
- Zhang, Z., Luo, P., Loy, C. C. and Tang, X. (2014) 'Facial landmark detection by deep multi-task learning', *ECCV*, pp. 94–108.
- University of Sussex (2026) *Applied Machine Learning*, lecture slides Wo6_L12 (image preprocessing), Wo9_L17 (HOG features), Wo9_L18 (landmark regression and augmentation), W10_L20 (data augmentation), W11_L21 (encoder–decoder architectures).